The Completion Fallacy in LLM Evaluation: A Process-Sensitive Probe for Cumulative State Tracking Predicts Agent Performance Beyond Task Completion

Dengzhe Hou[1,†] Lingyu Jiang[1] Deng Li[2] Zirui Li[1]
Fangzhou Lin[1,3,4] Michael Zielewski[1] Kazunori D Yamada[1]

[1]Tohoku University [2]Lappeenranta-Lahti University of Technology LUT
[3]Texas A&M University [4]Worcester Polytechnic Institute

hou.dengzhe.t8@dc.tohoku.ac.jp

[†]Corresponding author.

Abstract

LLM benchmarks overwhelmingly measure task completion but not how the model processes information—a gap we call the Completion Fallacy. We propose the Cognitive Evaluation Framework (CEF), a provisional research agenda of process-sensitive behavioral probes grounded in cognitive science, and report a construct-development pilot for one probe family, WMF-AM (cumulative arithmetic state tracking under load), as an existence proof that process-sensitive measurement adds value beyond completion. CEF's four-dimension structure is an unvalidated working taxonomy; WMF-AM is the sole thoroughly piloted sub-dimension.

In a 20-model open-weight pilot (13 families, 0.5B–35B parameters), WMF-AM significantly predicts performance on a deterministic 10-task multi-step agent battery (τ=0.612, p<0.001) and this signal persists after partialling out completion (τ_{partial}=0.411, p=0.011): a process probe captures downstream-relevant variance that a 100-item completion battery misses. Crucially, WMF-AM predicts even the 7 non-state-tracking agent tasks (τ=0.623, p<0.001), arguing against the simplest criterion-overlap explanation. A yoked cancellation control and non-arithmetic ablation show that WMF-AM difficulty is not explained by generic non-arithmetic state tracking, consistent with an arithmetic accumulation bottleneck; a completed last-operation-only (K=1) control confirms this: 13/15 models achieve near-ceiling single-step arithmetic accuracy ($\bar{x}$=0.957) yet floor at K=7 accumulation ($\bar{x}$=0.151), ruling out single-step arithmetic as the primary driver (though broader multi-step serial-composition alternatives cannot be excluded without further mechanistic analysis). All models are open-weight; generalization to cloud API models and robustness-oriented downstream tasks remains open.

1 Introduction

Current evaluation measures the products of cognition, not cognition itself. Benchmarks such as HotpotQA [Yang et al., 2018], MMLU [Hendrycks et al., 2021a], and AgentBench [Liu et al., 2023] tell us what LLMs accomplish, not how they process information. We call this the Completion Fallacy: the inference that task-completion success alone is sufficient for

inferring latent process competence. Completion metrics are often useful for outcome prediction, but completion alone does not distinguish agents that are reliable under distribution shift and process-coherent from those that are not—a special case of Goodhart's law [Goodhart, 1984, Manheim and Garrabrant, 2018].

Several recent results converge on this point. PAE [Xie et al., 2026] finds 27–78% of agent successes involve corrupted trajectories; Fragile Thoughts [Aravindan and Kejriwal, 2026] and Robust Answers Fragile Logic [Jiang et al., 2026] show correct answers produced despite broken reasoning chains; Turpin et al. [Turpin et al., 2023] demonstrate unfaithful chain-of-thought explanations. These are instances of multiple realizability [Putnam, 1967]: identical outcomes produced by categorically different processes—a pattern also documented as underspecification [D'Amour et al., 2022] and shortcut learning [Geirhos et al., 2020].

This paper is a position paper with a pilot study: we make a conceptual argument (Completion Fallacy), propose a provisional organizing framework (CEF), and provide empirical existence-proof evidence for one specific probe (WMF-AM). CEF's four-dimension structure is unvalidated; our evidence speaks to one sub-dimension.

This paper makes three contributions. (1) Formalization. We characterize the Completion Fallacy (Section 2) and identify three structural failure modes of outcome-only evaluation: one empirically illustrated (diagnostic blindness, with quantified model-pair evidence) and two hypothesized (uninformative failure diagnosis, safety blindness), each requiring further empirical confirmation. (2) Framework and pilot. We propose the Cognitive Evaluation Framework (CEF, Section 3), a provisional, unvalidated research agenda organized around four cognitive-science dimensions. A 20-model study (Section 4) provides existence-proof evidence that WMF-AM—a probe of cumulative arithmetic state tracking under load—captures process variance invisible to completion and predicts performance on a deterministic 10-task multi-step agent battery beyond completion (τ=0.612, p<0.001; partial τ=0.411, p=0.011), including for non-state-tracking tasks (τ=0.623). (3) Protocols. We provide complete experimental protocols for all 11 CEF sub-dimensions as a community resource (Appendix P). A glossary of all acronyms appears in Table 10 (Appendix C).

Three separable claims. (a) The general thesis: completion metrics alone are insufficient for inferring process quality. This claim is supported by converging independent evidence (PAE, Fragile Thoughts, Turpin et al.) and does not depend on CEF. (b) CEF as one proposed framework: a taxonomy of process-sensitive probes organized around cognitive science. Other frameworks are possible; CEF is offered as a concrete research agenda. (c) WMF-AM as one operationalization: the pilot provides existence-proof evidence for (a) through one specific probe family. The completion-fallacy claim does not rise or fall with any single sub-benchmark.

Evidence map. WMF-AM has multi-seed (all 15 models, 4 seeds), multi-template (3 templates, τ=0.524), yoked-control evidence, and downstream predictive validity (τ=0.612, N=20) (study-supported); EMC-lite, MCC-MA, and CLA-CR have single-seed evidence (preliminary); MCC-CE-v2 has initial results but zero variance (suggestive); remaining sub-dimensions have protocols only (proposed). Table 8 in Appendix B gives the full development status.

Scope of evidence. The study supports five claims: (i) process-sensitive probes can differentiate models that completion cannot (N=15, $\tau = 0.279$, bootstrap 95% CI $[-0.156, 0.637]$; 30/105 dissociation pairs), (ii) WMF-AM is reliable across seeds (4 per model), templates ($\tau_{\text{bare,chat}} = 0.524$, $p = 0.006$), and prompt paraphrases (mean cross-template $\tau = 0.54$; original-vs-formal $\tau = 0.857$), (iii) arithmetic accumulation—not generic entity tracking—drives WMF-AM difficulty (near-ceiling non-arithmetic ablation, mean 0.98), (iv) WMF-AM predicts performance on a 10-task deterministic multi-step agent battery beyond completion ($\tau = 0.612$, $p < 0.001$, N=20; partial $\tau = 0.411$, $p = 0.011$), and (v) the four-dimension CEF structure motivates further process-sensitive evaluation research; current data do not validate its dimensional structure, which awaits factor analysis at $N \geq 30$. The study does not support population-level generalization, stable dimensional separation, or robustness-

oriented downstream prediction. Sample size (N=20) supports basic inferential statistics but not factor analysis [Card et al., 2020].

Paper roadmap. The broad thesis—that completion metrics alone are insufficient for inferring process quality—is an independent conceptual claim supported by converging external evidence (Section 2). CEF is offered as a provisional organizing framework for process-sensitive probes (Section 3); it is a research agenda, not a validated instrument. WMF-AM is the only fully piloted sub-dimension: Section 4 reports its construct-development pilot evidence and downstream predictive validity as a concrete existence proof. All other dimensions (MCC, EMC, CLA) and their sub-probes are presented as hypotheses for future investigation, supported at most by single-seed preliminary results.

2 The Completion Fallacy and Related Work

Three failure modes (one empirically illustrated, two hypothesized). (F1) Diagnostic blindness [empirically illustrated]: models with near-identical completion scores can have very different process profiles. In the pilot study: 30/105 model pairs with $|\Delta\text{completion}| \leq 0.05$ show $|\Delta\text{WMF-AM}| \geq 0.10$ (exploratory; threshold-dependent—at $\leq 0.03/ \geq 0.15$, 18/105 pairs qualify). WMF-AM spans 0.050–0.983 among 20 models while completion spans only 0.44–0.92; among mid-to-large models (≥3B), completion SD = 0.052 versus WMF-AM SD = 0.228. (F2) Uninformative failure diagnosis [hypothesized]: current benchmarks report that an agent failed, not why—whether the bottleneck is state tracking, metacognitive calibration, or episodic confusion. Process-sensitive profiles, if validated at scale, could in principle decompose failure into actionable sub-dimensions. This function of CEF awaits full validation. (F3) Safety blindness [hypothesized]: pattern-matching agents (which fail predictably at distribution shift) are indistinguishable from genuine reasoners on standard benchmarks. Whether WMF-AM or other process probes predict this failure mode remains an open empirical question requiring dedicated robustness study.

Empirical support. Multiple realizability is empirically the norm in frontier LLMs. Fragile Thoughts [Aravindan and Kejriwal, 2026] and Robust Answers Fragile Logic [Jiang et al., 2026] provide direct evidence: correct answers are regularly produced despite corrupted reasoning steps. Turpin et al. [Turpin et al., 2023] show chain-of-thought explanations can be systematically unfaithful; Lanham et al. [Lanham et al., 2023] measure this faithfulness gap directly. Uesato et al. [Uesato et al., 2022] demonstrate that process-based feedback outperforms outcome-based feedback for mathematical reasoning. PAE [Xie et al., 2026] extends this to agent trajectories, finding 27–78% of successful completions involve corrupted intermediate actions. These are Agent B/C behaviors at scale.

From outcome underdetermination to process probes. Outcome metrics cannot distinguish process quality. Three questions follow: (a) What can task success not identify? Whether behavior generalizes under shift, whether confidence reflects error detection, or whether memory discriminates episodes. (b) Which gaps require process evidence versus robustness or adversarial testing? Process probes are needed when the question is how the model processes information. (c) Why CEF? Robustness testing [Ribeiro et al., 2020] addresses perturbation; trace audits [Turpin et al., 2023, Lanham et al., 2023] address reasoning-to-output coupling; adversarial evaluations address worst-case performance. CEF targets the space left uncovered: parametrically controlled, cognitively grounded probes of process quality under multi-turn load.

Positioning. Cognitive science has developed rigorous measurement paradigms for each function CEF addresses: working memory capacity [Miller, 1956, Baddeley and Hitch, 1974], metacognitive monitoring/control [Flavell, 1979, Nelson and Narens, 1990], episodic memory [Tulving et al., 1972, Johnson et al., 1993], and cognitive load adaptation [Sweller, 1988, Miyake et al., 2000]. CEF adapts their measurement logic for API-level LLM evaluation. This does not require LLMs to have human cognition; it requires that the constructs pick out capabilities that matter for deployment. NeuroCognition [Haznitrama et al., 2026] measures cognitive ability; CEF measures process quality under load. Strong Memory, Weak

Control [de Langis et al., 2025] studies executive function in isolated tasks; CEF tests the same constructs under sustained interference. Entity Tracking [Kim and Schuster, 2023] is prior art for WMF-AM; Intelligence Degradation [Wang et al., 2026] is prior art for CLA. Bowman and Dahl [Bowman and Dahl, 2021] argue benchmarking requires construct-level validity; CheckList [Ribeiro et al., 2020] introduced behavioral testing templates; CEF extends both. Process supervision [Lightman et al., 2023, Uesato et al., 2022] trains models using step-level labels; CEF measures process quality but does not prescribe training signals. Mechanistic interpretability analyzes internal representations; CEF operates at the API level (prompt/response only) and is complementary: interpretability explains why a process score is low, while CEF identifies that it is low. A full related-work table appears in Appendix A.

Formal definitions. Let $Y_i \in [0,1]$ denote the outcome score and $Z_i \in [0,1]$ a process-sensitive measure for model i, both defined relative to a specific benchmark suite and model population. We say outcome metrics are statistically insufficient with respect to a downstream criterion D if process measures provide significant incremental prediction of D beyond completion: partial $\tau(Z, D \mid Y) > 0$ at conventional α. This is a benchmark-conditional claim—it asserts insufficiency for a given $\langle Y, Z, D\rangle$ triple, not a universal property. Statistical insufficiency does not by itself determine practical or deployment relevance: the latter depends on whether the incremental signal is large enough and the criterion important enough to warrant the additional measurement cost. In the study, partial $\tau(\text{WMF-AM}, \text{Agent} \mid Y) = 0.411$ ($p = 0.011$, N=20), meeting this criterion. The completion–process correlation ($\tau_{\text{WMF-AM, completion}}$=0.279, p=0.150, N=15, bootstrap 95% CI $[-0.156, 0.637]$) is suggestive of divergence but not confirmed at this sample size. Falsification: CEF's core claim would be undermined if, across multiple probe families and downstream criteria, process measures never provided significant incremental prediction beyond completion. The four-dimension structure is a working taxonomy; dimensional separability requires factor analysis at $N \geq 30$.

Statistical methodology. All rank correlations use Kendall's τ-b (accounts for ties). Partial τ uses Somers residualization:

$$\tau_{XY|Z} = \frac{\tau_{XY} - \tau_{XZ}\,\tau_{YZ}}{\sqrt{(1-\tau_{XZ}^2)(1-\tau_{YZ}^2)}}.$$

Bootstrap 95% CIs: 10,000 nonparametric model-level resamples (percentile method). Permutation tests: 100,000 permutations. All p-values are two-sided. The two pre-specified confirmatory analyses are predictive validity ($\tau = 0.612$) and template stability ($\tau = 0.524$); both survive Bonferroni correction at $\alpha = 0.025$. All other correlations are exploratory.

Table 1. Key analyses at a glance. All τ = Kendall's τ-b. C = confirmatory (pre-specified, Bonferroni-corrected); E = exploratory.

Relation	Subset	τ	N	p	95% CI	Type
WMF-AM → Agent overall	All models	0.612	20	<0.001	$[0.360, 0.814]$	C
Bare ↔ Chat template ranks	15 models	0.524	15	0.006	—	C
WMF-AM → Agent (partial | completion)	All models	0.411	20	0.011	—	E
WMF-AM → Agent (partial | log params)	All models	0.503	20	<0.01	—	E
WMF-AM → Agent (partial | compl. + log params)	All models	0.389	20	0.016	—	E
WMF-AM → Agent non-WMF (7 tasks)	All models	0.623	20	<0.001	—	E
WMF-AM → Agent WMF-heavy (3 tasks)	All models	0.527	20	0.003	—	E
WMF-AM ↔ completion	15 models	0.279	15	0.150	$[-0.156, 0.637]$	E
CONV-RFPOC ↔ AGENT-PQ	7 models	0.905	7	0.003	—	E

3 The Cognitive Evaluation Framework (CEF)

CEF is a hypothesis-generating taxonomy—an organizational framework for process-sensitive probes, not an empirically validated instrument. We emphasize this hierarchy explicitly: the paper presents one thoroughly piloted probe family (WMF-AM) as a concrete existence

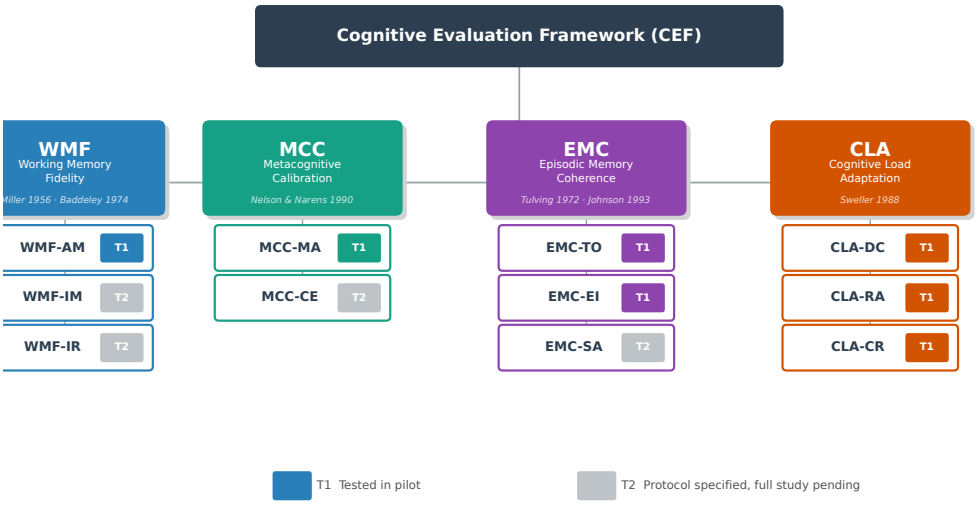

Figure 1. CEF architecture. Four dimensions (WMF, MCC, EMC, CLA) derived from cognitive science, decomposed into 11 proposed sub-dimensions. Solid border: pilot-supported (WMF-AM). Dashed border: preliminary or proposed probes awaiting the $N \geq 15$ study.

proof, embedded within a broader speculative taxonomy (CEF) that offers candidate dimensions for future investigation. The taxonomy should be read as a research agenda, not as a validated instrument. Only WMF-AM has multi-evidence pilot support (predictive validity, seed reliability, template stability, yoked control, arithmetic-vs-assignment ablation, paraphrase stability); three sub-dimensions have preliminary single-seed results; the remainder are protocol-only (Appendix B, Table 8). The taxonomy is offered to organize future research; the number and boundaries of dimensions are working hypotheses pending factor analysis at $N \geq 30$.

Design principles. Each dimension must have (i) a cognitive science anchor—an established measurement paradigm with validated human operationalizations, (ii) a documented evaluation gap not covered by existing benchmarks, and (iii) API-level testability requiring only prompt/response access. We use labels like "working memory" as operational shorthand for behavioral probes, not as claims about LLM mechanisms. "WMF-AM" means accuracy on sequential state-tracking tasks degrades with depth K; it does not mean LLMs have a working memory buffer (see Appendix D for full discussion).

Four dimensions. Table 2 summarizes each dimension's target capacity, canonical failure signature, and exclusions.

Flagship probe: WMF-AM. WMF-AM (Active Manipulation) is a probe of cumulative arithmetic state tracking under load, inspired by working-memory span paradigms. An initial numeric state is presented, K sequential additive/subtractive operations arrive, and the model must report the final cumulative state. The pilot uses depths $K \in \{3, 5, 7\}$ (dissociation study) and $K \in \{2, 4, 6, 8, 12\}$ (yoked cancellation control), across three surface forms (points scoring, warehouse inventory, bank accounts). Prior entity tracking [Kim and Schuster, 2023] tests passive retrieval; WMF-AM tests sequential arithmetic transformation under load. Scope of the construct: WMF-AM specifically measures difficulty with maintaining a running cumulative sum across K additive/subtractive operations. It does not claim to measure general working-memory capacity, and it may partially reflect standalone arithmetic competence; a completed last-operation-only (K=1) control addresses this: 13/15 models achieve near-ceiling single-step arithmetic accuracy ($\bar{x}$=0.957) versus $\bar{x}$=0.151 at K=7, indicating that cumulative accumulation under load—not arithmetic skill per se—is the primary difficulty driver.

Worked example (K=3, points form): "Alice starts with 10 points. Alice gains 5 points.

Table 2. CEF dimensions: target, failure signature, and exclusions. Detailed sub-dimension definitions appear in Appendix E.

Dim	Target	Failure signature	Excluded
WMF	Active maintenance and transformation of state under load	Accuracy drops as concurrent items exceed threshold	Long-context retrieval; knowledge recall
MCC	Self-monitoring of errors (MA) and correction of flagged errors (CE)	Confident-wrong; flagging without correction	General calibration (ECE); cross-model detection
EMC	Temporal ordering, cross-episode discrimination, source attribution	Episode blending; source misattribution	Single-turn retrieval; semantic memory
CLA	Degradation shape and resource scaling under load	Cliff-edge failure; non-recovery	Knowledge difficulty; unrelated complexity

Table 3. Metacognitive profile taxonomy. Each cell implies a different intervention target.

	High MCC-CE	Low MCC-CE
High MCC-MA	Genuine metacognition	Monitoring works; correction broken
Low MCC-MA	Correction works; detection unreliable	No metacognitive capability

Alice loses 3 points. Alice gains 7 points. What is Alice's current score?" Correct answer: 19. Scoring: exact match (1 if the model outputs 19, 0 otherwise). A model that tracks each update cumulatively ($10 \rightarrow 15 \rightarrow 12 \rightarrow 19$) succeeds; a model that performs only the last operation or guesses a plausible number fails. At K=7, the chain length taxes state maintenance, producing the diagnostic depth×accuracy interaction. Yoked control example (K=3): "Alice starts with 10 points. Alice gains 5 points. Alice loses 5 points. Alice gains 3 points. Alice loses 3 points. Alice gains 7 points. Alice loses 7 points. What is Alice's current score?" Correct answer: 10 (initial state). Each operation is immediately followed by its inverse, requiring identical arithmetic but no cumulative tracking.

MCC: Metacognitive Calibration. MCC decomposes metacognition into monitoring accuracy (MCC-MA: predict which of 20 answers are wrong; Pearson r) and control efficacy (MCC-CE: flag uncertain answers, then revise them; decomposed into flagging rate, correction efficacy, and false alarm rate)—Reflexion [Shinn et al., 2023] conflates these. Unlike Gnosis [Ghasemabadi and Niu, 2025] (requires internals), MCC-MA uses prompting under multi-turn load. The 2×2 taxonomy (Table 3) makes the decomposition actionable: each cell implies a different architectural intervention.

EMC: Episodic Memory Coherence. EMC tests temporal ordering (EMC-TO), cross-episode discrimination (EMC-EI), and source attribution (EMC-SA) under interference. Current benchmarks measure binary retrieval; no protocol tests whether a model situates retrieved facts in time and context. The EMC-lite pilot (Section 4) covers a related but narrower slice.

CLA: Cognitive Load Adaptation. CLA characterizes the shape of the performance-versus-load function. CLA-DC (degradation curve smoothness) correlates highly with WMF-AM ($\tau = 0.905$) in the original 7-model set. We provisionally collapse CLA-DC onto WMF: both involve state tracking under load. Factor analysis with N=15 models is planned to determine whether this collapse holds. This does not invalidate the taxonomy—it refines it. CLA-RA (resource allocation) and CLA-CR (capacity recovery after high-load block) are structurally

Table 4. Study design summary. All phases: temperature = 0, greedy decoding, Ollama-hosted models with default chat templates. All 20 models evaluated with 4 seeds for WMF-AM.

Phase	Probe	Depths	Items	Seeds	Scoring
1. Outcome	100-item battery (5 domains)	—	100	1	Exact match (0/1)
2. WMF-AM	State tracking	K=3, 5, 7	15/depth	4	Exact match on final state
3. MCC-MA	Self-monitoring	—	20	1	Pearson r (predicted vs actual)
4. Yoked	Cancellation ops	K=2, 4, 6, 8, 12	20/depth	1	Exact match on initial state
5. Template	3 prompt wrappers	K=3, 5, 7	15/depth	1	Rank stability (τ)
6. Agent	10 multi-step tasks	—	10	1	Process quality (0–1)

independent and cannot collapse onto WMF-AM by design (they measure recovery, not degradation). Intelligence Degradation [Wang et al., 2026] reports that performance drops; CLA characterizes how. Full definitions, composites, and scoring rules are in Appendix E; Table 9 in Appendix B summarizes all 11 sub-dimensions.

Construct validity framework. Following Cronbach and Meehl [Cronbach and Meehl, 1955], Messick [Messick, 1989], Campbell and Fiske [Campbell and Fiske, 1959], and Jacobs and Wallach [Jacobs and Wallach, 2021], each dimension requires convergent validity (correlation with independent process-sensitive measures) and divergent validity (low correlation with completion). The study supports divergent validity for WMF-AM ($\tau = 0.279$ with a 100-item completion battery, N=15; bootstrap 95% CI $[-0.156, 0.637]$) and cross-template rank stability ($\tau_{\text{bare,chat}} = 0.524$, $p = 0.006$). CEF would be falsified if all dimensions collapsed onto a single factor indistinguishable from task completion. Appendix F provides full construct-to-probe mapping with confounds and falsification criteria for each pilot-tested probe.

4 Pilot Study: Process-Sensitive Differentiation

Note on sample sizes. Different analyses use different subsets: N=20 for all predictive validity and completion-battery analyses (the full study set); N=15 for WMF-AM template stability and divergent validity (original 7 + expansion 8; before small models were added); N=7 for convergent validity only (original models with AGENT-PQ and CONV-RFPOC data collected). All sample sizes are reported explicitly at each result.

Twenty open-weight models from thirteen architectural families form the study set: the original seven (Qwen 7B/14B/32B, Llama 3.1:8b, Gemma 2:27b, DeepSeek-R1:14b, Mistral:7b), eight expansion models (Phi-3:14b, Gemma-2:9b, Qwen-2.5:3b, Llama-3.2:3b, DeepSeek-R1:7b, Mixtral:8x7b, Command-R:35b, Yi:34b), and five small models for broader baseline coverage (Qwen-2.5:0.5b/1.5b, Llama-3.2:1b, Gemma-2:2b, TinyLlama:1.1b; 0.5B–2B parameters). Models were evaluated in six phases. Phase 1 (Outcome Correctness) used a 100-item battery spanning five domains (factual, math, science, reasoning, language) at five difficulty levels, temperature = 0. Phase 2 (WMF-AM) administered 15 probes at depths $K \in \{3, 5, 7\}$ with 4 independent random seeds for all 20 models. Phase 3 (MCC-MA) applied a 20-problem self-monitoring probe. Phase 4 (Yoked Cancellation Control) administered 20 trials per depth at $K \in \{2, 4, 6, 8, 12\}$ using operations that cancel (gain then lose, transfer then reverse), so the correct answer equals the initial state despite requiring identical arithmetic parsing as WMF-AM. Phase 5 (Template Harmonization) re-administered WMF-AM under three prompt wrappers (bare, chat, chain-of-thought) to test sensitivity to surface format. Phase 6 (Agent Validation) administered 10 multi-step agent tasks (tool use, multi-step reasoning, state tracking) to all 20 models to test predictive validity. All models are Ollama-hosted open-weight variants using default chat templates, fp16 quantization where available, and greedy decoding (temperature = 0, no top-p sampling). No task overlap exists between phases (full study matrix in Appendix H). Table 4 summarizes the experimental design.

Main result: process measures provide incremental prediction beyond completion on a multi-step agent battery. Figure 2 and Table 5 present the key finding across N=20 models from 13 families. WMF-AM (4-seed means) spans 0.050–0.983 ($\Delta = 0.933$) while the 100-item

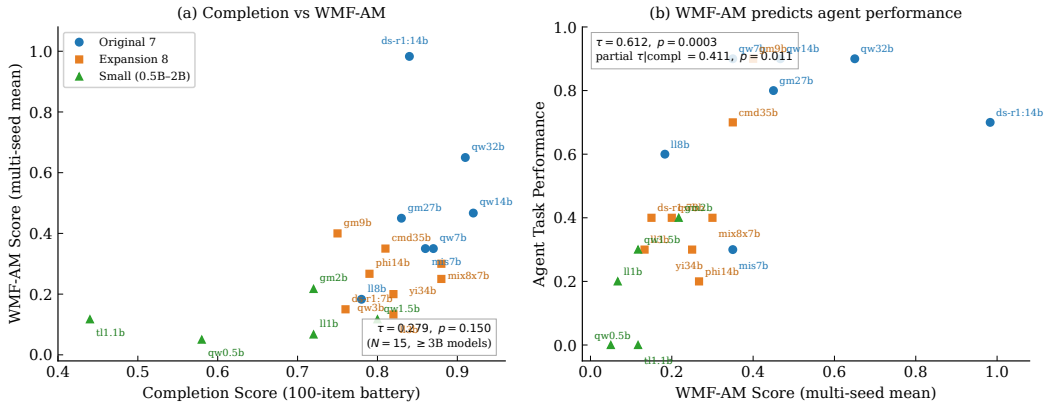

Figure 2. (a) Completion clusters while WMF-AM spans a wide range (N=20, 13 families). (b) WMF-AM predicts downstream agent performance (τ=0.612, p<0.001) even after controlling for completion (partial τ=0.411, p=0.011). Blue circles = original 7; orange squares = expansion 8; green triangles = small models (0.5B–2B).

completion battery spans 0.44–0.92 ($\Delta = 0.480$). Among the 15 mid-to-large models ($\geq$3B), completion clusters at 0.75–0.92 (SD = 0.052) while WMF-AM spans 0.133–0.983 (SD = 0.228): completion masks substantial process-quality differences. Predictive validity (scoped claim): WMF-AM significantly predicts performance on a deterministic 10-task multi-step agent battery ($\tau = 0.612$, $p < 0.001$, N=20; bootstrap 95% CI $[0.360, 0.814]$). Crucially, this relationship holds after partialling out completion: partial $\tau = 0.411$ ($p = 0.011$), demonstrating that WMF-AM captures information about downstream capability beyond what completion provides. Completion also predicts agent performance ($\tau = 0.428$, $p = 0.012$), but WMF-AM provides incremental prediction. The primary evidence is threefold: (1) significant predictive validity, (2) significant template stability ($\tau = 0.524$, $p = 0.006$), and (3) WMF-AM reliably differentiates models in a dimension that completion does not fully capture.

Completion battery design and broader baseline. Phase 1 uses a 100-item battery spanning five domains (factual, math, science, reasoning, language) at five difficulty levels, yielding score resolution of 0.01. With the addition of 5 small models (0.5B–2B), scores now range from 0.44 (tinyllama:1.1b) to 0.92 (qwen2.5:14b) (SD = 0.131, N=20), a substantial improvement over the original 15-model range of 0.75–0.92 (SD = 0.052). The broader range strengthens interpretation: completion now correlates with agent performance ($\tau = 0.428$, $p = 0.012$), but WMF-AM provides significant incremental prediction (partial $\tau = 0.411$, $p = 0.011$). Among the 15 mid-to-large models ($\geq$3B), the outcome range remains compressed (SD = 0.052), confirming that range restriction is a property of the model population, not the battery design. Difficulty-level analysis confirms the battery discriminates (level-5 items: 0.0–1.0 spread; level-1 items: 0.73–0.95).

Exploratory architectural observations. Reasoning-specialized architecture (deepseek-r1:14b, 0.983) outperforms general-purpose models despite similar parameter counts. Scaling within families yields modest gradients (Qwen 3B→32B: 0.311 → 0.650). The yoked control reveals an important confound decomposition: only deepseek-r1:14b shows positive WMF−Yoked residual (+0.173), suggesting it alone achieves genuine cumulative state tracking beyond per-operation parsing. All other models show negative residuals, with the largest gaps in smaller models (deepseek-r1:7b: −0.829). Figure 4 shows the depth-dependent degradation curves for all 15 models.

Alternative explanations and template sensitivity. WMF-AM scores may reflect nuisance factors beyond state tracking: instruction-following ability, tokenizer effects on numeric parsing, scratchpad/chain-of-thought policy differences, or prompt-format sensitivity. Phase 5 (template harmonization) directly tests the last concern: across 15 models × 3 templates,

Model	Family	Outcome	WMF-AM	Agent	Yoked	WMF−Yoked
Original 7 (all 4-seed WMF-AM)						
deepseek-r1:14b	DeepSeek	0.840	0.983	0.70	0.810	+0.173
qwen2.5:32b	Qwen	0.910	0.650	0.90	1.000	−0.350
qwen2.5:14b	Qwen	0.920	0.467	0.90	0.710	−0.243
gemma2:27b	Gemma	0.830	0.450	0.80	0.620	−0.170
qwen2.5:7b	Qwen	0.870	0.350	0.90	0.770	−0.420
mistral:7b	Mistral	0.860	0.350	0.30	0.550	−0.200
llama3.1:8b	Llama	0.780	0.183	0.60	0.340	−0.157
Expansion 8 (all 4-seed WMF-AM)						
gemma2:9b	Gemma	0.750	0.400	0.90	0.960	−0.560
command-r:35b	Cohere	0.810	0.350	0.70	0.920	−0.570
mixtral:8x7b	Mistral	0.880	0.300	0.40	1.000	−0.700
phi3:14b	Phi	0.790	0.267	0.20	0.950	−0.683
yi:34b	Yi	0.880	0.250	0.30	1.000	−0.750
qwen2.5:3b	Qwen	0.820	0.200	0.40	0.970	−0.770
deepseek-r1:7b	DeepSeek	0.760	0.150	0.40	0.940	−0.790
llama3.2:3b	Llama	0.820	0.133	0.30	0.510	−0.377
Small models (broader baseline)						
gemma2:2b	Gemma	0.720	0.217	0.40	—	—
qwen2.5:1.5b	Qwen	0.800	0.117	0.30	—	—
tinyllama:1.1b	TinyLlama	0.440	0.117	0.00	—	—
llama3.2:1b	Llama	0.720	0.067	0.20	—	—
qwen2.5:0.5b	Qwen	0.580	0.050	0.00	—	—
Δ	—	0.480	0.933	0.90	—	—

Table 5. Pilot study (N=20, 13 families). Outcome = 100-item battery; WMF-AM = 4-seed mean; Agent = multi-step task performance (10 tasks). WMF-AM predicts agent performance ($\tau = 0.612$, $p < 0.001$) even after controlling for completion (partial $\tau = 0.411$, $p = 0.011$). Yoked column shows cancellation-control accuracy (original 15 only).

model rankings under bare and chat wrappers are significantly correlated ($\tau_{\text{bare,chat}} = 0.524$, $p = 0.006$), indicating that WMF-AM rankings are moderately robust to surface-format variation. Chain-of-thought (CoT) templates push nearly all models to ceiling (≥ 0.93), eliminating variance.

Target construct clarification: WMF-AM is designed to measure default state-tracking capability under a fixed prompt regime (no scratchpad, no chain-of-thought). It deliberately probes the model's internalized processing rather than its externalized reasoning when given scaffolding. The CoT ceiling effect is therefore not a validity threat but a construct boundary: WMF-AM measures what the model can do without externalized reasoning aids [Nye et al., 2021], which is the relevant capability for deployment contexts where intermediate outputs are not monitored. The moderate bare↔chat stability ($\tau = 0.524$, not near 1.0) indicates some prompt sensitivity remains; this is a known limitation. The yoked control isolates arithmetic parsing, but instruction-following and tokenizer effects remain live confounds. Ablation results (N=15): (i) Non-arithmetic state tracking (3 domains $\times$ $K \in \{3, 5, 7\}$, 10 problems each): when operations involve direct-assignment updates (color "Alice's color changes to blue", location, status) instead of numeric accumulation, models achieve near-ceiling accuracy (mean 0.98, range 0.91–1.00), versus the WMF-AM arithmetic range of 0.20–0.72. Direct-assignment tracking is solved by remembering the most recent value per entity; arithmetic tracking requires maintaining a running cumulative sum across K additive/subtractive operations—a strictly harder representational demand. The floor–ceiling contrast confirms that WMF-AM difficulty is driven by arithmetic accumulation under load, not generic entity-state tracking. This does not rule out that WMF-AM partially reflects standalone arithmetic skill; a completed last-operation-

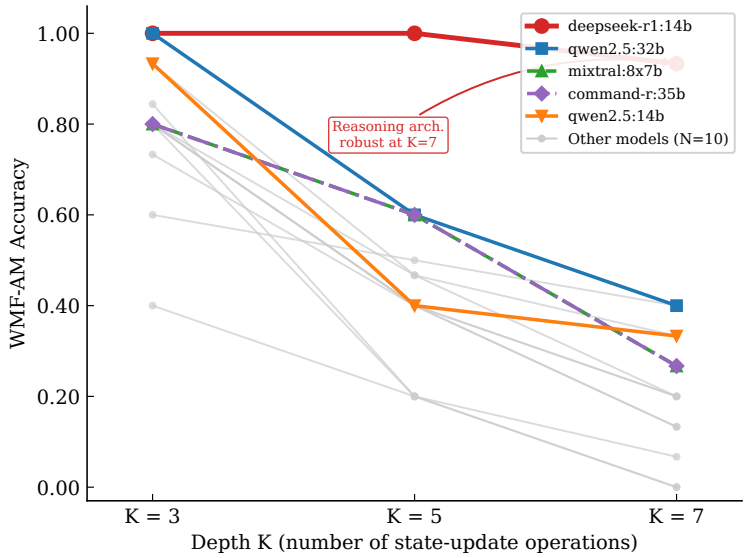

Figure 3. WMF-AM accuracy by depth K (N=15, chat template). deepseek-r1:14b (red) maintains near-perfect accuracy; all others degrade sharply at $K \geq 5$. Top 5 models highlighted.

only (K=1) control addresses this: 13/15 models achieve near-ceiling single-step arithmetic accuracy (mean 0.957) yet floor at K=7 accumulation (mean 0.151), ruling out single-step arithmetic as the primary difficulty driver. Broader multi-step serial-composition alternatives (e.g., error propagation, sequential prompt processing) cannot be excluded without further mechanistic analysis. (ii) Prompt paraphrase (5 templates: original, formal, casual, minimal, verbose): model rankings are stable across natural-language rephrasings; mean cross-template Kendall $\tau = 0.54$ (5/6 pairs $p < 0.05$); original-vs-formal $\tau = 0.857$ ($p < 0.001$). The stripped "minimal" template induces task failure rather than rank reversal, consistent with insufficient-context degradation rather than construct invalidity. Rankings on the paraphrase-original template correlate near-perfectly with the main WMF-AM results: $\tau = 0.951$ ($p = 0.004$, N=7 overlapping models). (iii) Completed in prior loops: three surface forms (points, inventory, accounts), three prompt wrappers (bare, chat, CoT), and four random seeds—all showing stable inter-model rank orderings.

Predictive validity: agent task performance. Phase 6 administered 10 multi-step agent tasks (3 requiring tool use with state tracking, 3 multi-step reasoning chains, 2 general problem-solving, 2 constrained planning) to all 20 models. Each task is scored 0–1 by deterministic verification: intermediate and final states are checked against ground-truth answers via exact match (no LLM judge). This yields the Agent score used in all predictive validity analyses. (A separate LLM-judged process quality metric, AGENT-PQ, is available for the original 7 models only and used solely in convergent validity; see Appendix O.) Agent scores range from 0.00 (qwen2.5:0.5b, tinyllama:1.1b) to 0.90 (qwen2.5:14b/32b/7b, gemma2:9b).

WMF-AM significantly predicts agent performance: $\tau = 0.612$ ($p < 0.001$, N=20; bootstrap 95% CI $[0.360, 0.814]$). Confound controls: (i) Partial τ controlling for completion: 0.411 ($p = 0.011$). (ii) Partial τ controlling for log(parameter count): 0.503 ($p < 0.01$, exploratory)—the relationship is not driven by model scale. (iii-joint) Partial τ controlling for both completion and log(parameter count) simultaneously: 0.389 ($p = 0.016$, exploratory)—WMF-AM retains significant incremental signal after accounting for both outcome performance and model size jointly. (iii) Within-family analysis: in the Qwen family (6 models, 0.5B–32B), $\tau = 0.894$ ($p = 0.016$), demonstrating that WMF-AM predicts agent performance within a single architecture family, not just across families. Criterion contamination control (pre-registered split): To rule out the most direct objection—that WMF-AM predicts the agent battery merely because both involve state tracking—we split the 10 tasks

into WMF-heavy (3 tool-use tasks explicitly requiring sequential state tracking: inventory management, multi-step booking, resource allocation) and non-WMF (7 tasks: multi-step reasoning, planning, general problem-solving, constrained optimization—none require sequential cumulative state tracking). WMF-AM predicts non-WMF tasks more strongly ($\tau = 0.623$, $p < 0.001$) than WMF-heavy tasks ($\tau = 0.527$, $p = 0.003$), arguing against the simplest criterion-overlap explanation. Shared latent demands (sequential execution discipline, multistep consistency) cannot be fully excluded, but the split result undermines the most direct contamination account. Leave-one-out robustness: removing any single model, the WMF-AM $\rightarrow$ agent τ ranges $[0.537, 0.624]$ (mean 0.582)—no single model drives the effect. Scope: these results demonstrate incremental predictive signal beyond completion and scale in this sample of open-weight models. Generalization to cloud API models and robustness under distribution shift remain open questions.

Yoked cancellation control. To isolate active state tracking from per-operation arithmetic parsing, we designed a stricter control: each operation is immediately followed by its inverse ("Alice gains 3 points" then "Alice loses 3 points"), so the correct answer always equals the initial state despite requiring identical arithmetic processing as WMF-AM. This eliminates the shortcut in the original inert control (where the answer was trivially the initial state without any computation).

Result: Table 5 (Yoked column) shows that most models achieve high yoked-control accuracy (0.51–1.00) but substantially lower WMF-AM accuracy, indicating the WMF-AM difficulty is genuine cumulative state tracking beyond arithmetic parsing. The corrected residual (WMF−Yoked) is negative for 14/15 models, meaning per-operation parsing inflates raw WMF-AM scores. With the 100-item completion battery, τ(WMF−Yoked, Outcome) = 0.125 ($p = 0.519$, N=15) and raw WMF-AM $\tau = 0.279$ ($p = 0.150$): neither the raw nor the residual score significantly predicts outcome quality, strengthening the independence claim. The near-universal negative WMF−Yoked residuals (14/15 models) confirm that per-operation arithmetic parsing substantially inflates raw WMF-AM scores, and the yoked control successfully decomposes WMF-AM into a state-tracking component and an arithmetic-parsing confound.

Measurement robustness. Multi-seed reliability: all 15 models evaluated with 4 independent seeds $\times$ 15 probes. Expansion-model SDs range 0.073–0.144 (mean 0.108), comparable to original-model variability. Cross-template stability: $\tau_{\text{bare,chat}} = 0.524$ ($p = 0.006$); $\tau_{\text{chat,cot}} = 0.314$ ($p = 0.114$). See Appendix J for full analyses.

Complementary probes (preliminary; see Appendix K–M). Three additional probes were run at pilot scale ($N \leq 8$, single seed) and should be treated as preliminary existence checks, not validated sub-dimensions. EMC-lite (a source-conflict episode attribution probe) shows spread (= 0.400) with low correlation to WMF-AM ($\tau = 0.250$), consistent with dimensional separability but insufficient to confirm it. MCC-CE-v2 (a control-efficacy probe: model flags uncertain answers and revises) eliminated the v1 floor effect (error rates 23–40%) but found 7/8 models with zero self-flagging, suggesting open-weight models may fundamentally lack error-correction metacognition under the binary protocol—a finding requiring replication. CLA-DC (degradation-curve smoothness over load levels) provisionally collapses onto WMF-AM ($\tau = 0.905$); pending factor analysis at $N \geq 20$. CLA-CR-v3 (capacity recovery after high-load block) shows spread but is two-seed only. We draw no multi-dimensional structure claims from these pilots; they motivate but do not validate the CEF taxonomy.

4.1 Validity Summary

Table 6 summarizes validity evidence. Predictive validity uses N=20; divergent statistics use N=15; convergent validity uses the original 7 models (the only set with AGENT-PQ/CONV-RFPOC data).

Predictive validity (Table 6, top block) is the strongest result: WMF-AM predicts agent performance on this open-weight deterministic battery beyond completion and scale, surviving all four confound controls (including joint partial) (Section 4). Convergent: CONV-

Table 6. Validity summary. Predictive: N=20; divergent: N=15; convergent: N=7 (original set). All τ = Kendall's τ-b. C = confirmatory (pre-specified, Bonferroni-corrected); E = exploratory.

Analysis	Pair	τ	p	95% CI	Type
Predictive (N=20)					
	WMF-AM → Agent overall	0.612	0.0003	[0.360, 0.814]	C
	WMF-AM → Agent (partial | completion)	0.411	0.011	—	E
	WMF-AM → Agent (partial | log params)	0.503	<0.01	—	E
	WMF-AM → Agent (partial | compl. + log params)	0.389	0.016	—	E
	WMF-AM → Agent (within-Qwen, N=6)	0.894	0.016	—	E
	WMF-AM → Agent WMF-heavy (3 tasks)	0.527	0.003	—	E
	WMF-AM → Agent non-WMF (7 tasks)	0.623	0.0002	—	E
	Completion → Agent overall	0.428	0.012	[0.031, 0.743]	E
Convergent (N=7)					
	CONV-RFPOC vs AGENT-PQ	0.905	0.003	—	E
	AGENT-TC vs AGENT-PQ	0.720	0.028	—	E
Divergent (N=15)					
	WMF-AM vs MCC-MA	0.169	0.615	—	E
	WMF-AM vs Completion (100-item battery)	0.279	0.150	[−0.156, 0.637]	E
	WMF−Yoked vs Outcome (100-item)	0.125	0.519	—	E
	Bare vs Chat template ranks	0.524	0.006	—	C
Cross-dimension (N=7)					
	WMF-AM vs CLA-DC	0.905	0.003	—	E
	WMF-AM vs EMC-lite	0.250	0.442	—	E

RFPOC (Reasoning Fidelity–Product of Cognition: the fraction of reasoning problems whose answer changes when the chain is counterfactually perturbed; higher = more faithful reasoning) ↔ AGENT-PQ (Agent Process Quality: LLM-judged rubric score on the original 7 models only, distinct from the deterministic Agent score above) yields $\tau = 0.905$ (N=7). Caveat: AGENT-PQ uses an LLM judge (GPT-4o); the high convergence may partly reflect shared evaluator bias. Human annotation is needed to confirm. Divergent: WMF-AM ↔ completion $\tau = 0.279$ (N=15, 100-item battery); WMF-AM ↔ MCC-MA $\tau = 0.169$. Neither raw WMF-AM nor the yoked-corrected residual (WMF−Yoked, $\tau = 0.125$, $p = 0.519$) significantly predicts outcome, supporting the independence of process measurement from task completion. Construct collapse: WMF-AM ↔ CLA-DC $\tau = 0.905$; provisionally collapsed pending $N \geq 20$. Even if CLA-DC merges with WMF-AM, WMF-AM, MCC-MA, and EMC-lite remain separable ($\tau \leq 0.250$). LOMO: convergent τ ranges 0.600–1.000 (mean 0.833). Template stability: bare ↔ chat template rankings yield $\tau = 0.524$ ($p = 0.006$), surviving Bonferroni correction. Full details in Appendix N.

5 Limitations, Alternative Views, and Discussion

Objection 1: Outcomes are what matter. We accept this for narrow, stable domains. For deployment settings with distribution shift and asymmetric failure costs, process quality matters. Predictive validity is now partially addressed: WMF-AM predicts performance on a deterministic 10-task multi-step agent battery in this open-weight pilot ($\tau = 0.612$, $p < 0.001$, N=20), and this holds after jointly controlling for completion and scale (partial $\tau = 0.389$, $p = 0.016$). This demonstrates that process-sensitive measurement captures downstream-relevant variance that completion misses, on this specific battery. Open question: does this extend to robustness under distribution shift (prompt perturbation, long-horizon failure, adversarial inputs)? The current 10 agent tasks test multi-step deterministic capability; establishing whether WMF-AM predicts robustness-oriented failure modes (corrupted intermediate states, distribution shift, adversarial inputs) remains the highest-priority item for the full study.

Objection 2: Cognitive paradigms were not designed for LLMs. CEF sub-dimensions are behavioral probes with well-defined stimuli, independent of mechanism. The question is whether probes measure something useful, not whether LLMs have cognition. The dissoci-

ation pilot demonstrates that WMF-AM reliably differentiates models in ways completion cannot.

Objection 3: Sample size. The current study uses N=20 models from 13 families (0.5B–35B), addressing the original N=7 limitation. While still insufficient for factor analysis [Card et al., 2020], the sample supports multiple significant inferential results: predictive validity ($\tau = 0.612$, $p < 0.001$), template stability ($\tau = 0.524$, $p = 0.006$), and partial correlation ($\tau = 0.411$, $p = 0.011$). Held-out evaluation with cloud API models is planned for external validity and broader parameter-range coverage.

Key limitations. (i) All 20 models are open-weight Ollama variants; held-out evaluation with cloud API models (GPT-4o, Claude, Gemini) is needed for external validity. (ii) Agent tasks test multi-step capability, not robustness under distribution shift; robustness-predictive validity remains open. (iii) WMF-AM is a probe of cumulative arithmetic state tracking; it does not measure general working-memory fidelity. A completed last-operation-only control (K=1, N=15 models, 90 trials each) addresses this: 13/15 models score near-ceiling on standalone arithmetic ($\bar{x}$=0.957) yet floor at K=7 ($\bar{x}$=0.151), indicating cumulative accumulation under load—not arithmetic competence per se—as the primary source of WMF-AM difficulty. (Two exceptions, deepseek-r1:14b and deepseek-r1:7b, show anomalously low K=1 exact-match scores attributable to a parsing artifact: their verbose narrative responses cause the first-number extractor to return the initial value rather than the final answer; 88/90 deepseek-r1:14b responses contain the arithmetically correct result.) (iv) MCC-CE-v2 uses a binary protocol that may under-elicit partial metacognitive capacity; preliminary evidence only. (v) The deterministic Agent score is the primary downstream criterion; AGENT-PQ (LLM-judged, N=7 only) appears only in convergent validity and requires human annotation to confirm. (vi) Rival explanations remain open for each probe (Appendix G). (vii) Template harmonization shows bare↔chat stability (τ=0.524), but CoT collapses all variance—WMF-AM measures default state-tracking capability in the constrained-externalization regime (no scratchpad). This is a construct boundary, not a failure. (viii) N=20 supports basic inference but not factor analysis; the four-dimension structure is a provisional organizing taxonomy, not an empirically validated model. (ix) Family non-independence: the 20-model set includes 6 Qwen models and 3 DeepSeek models; rank correlations may be inflated if within-family ordering is trivial. The within-Qwen analysis (τ=0.894, p=0.016, N=6) is included as a robustness check, but family-balanced sampling is needed to rule out this confound fully. (x) The K=1 control uses first-number extraction; two verbose reasoning models (deepseek-r1:14b/7b) show anomalously low exact-match scores due to this parser. Qualitatively, 88/90 deepseek-r1:14b responses contain the correct arithmetic result, and the conclusion (arithmetic is not the bottleneck for 13/15 models) is robust to this artifact; however, a parser-robust or manually annotated replication is needed for full rigor.

Separating the thesis from the framework. The completion-fallacy claim—that completion alone is insufficient for inferring process quality—is supported by converging independent evidence and does not depend on CEF or WMF-AM specifically. The empirical contribution is evidence for one process-sensitive probe family (WMF-AM), situated within a broader hypothesis-generating taxonomy (CEF). Other process-sensitive frameworks are possible and welcome. The pilot provides an existence proof, not a claim that WMF-AM is the uniquely correct construct or that CEF's four-dimension structure is empirically confirmed.

Implications for benchmark design. We suggest benchmarks would benefit from including at least one process measure alongside completion: (i) a probe testing the causal connection between reasoning and outcome [Aravindan and Kejriwal, 2026, Lightman et al., 2023], (ii) performance across a difficulty gradient rather than a single operating point, and (iii) a metacognitive monitoring probe enabling the monitoring/control dissociation to surface.

Implications for development (hypothesized). If supported at scale, CEF profiles could serve as diagnostic indicators: low WMF-AM points to weak state tracking; the MCC decomposition identifies whether the bottleneck is detection or correction—the pilot finding of near-zero self-flagging in open-weight models suggests self-correction architectures may require

external verification; cliff-edge CLA-DC signatures may predict distribution-shift fragility. These are testable hypotheses, not established claims.

Research agenda. Three directions merit community investigation, contingent on full study results ($N \geq 15$): (1) benchmark developers include load-varied conditions and metacognitive probes; (2) LLM developers report process-quality profiles alongside task performance; (3) the community investigates whether process validity complements task validity as a first-class evaluation concern. CEF is a proposed research program, not a finished instrument; the WMF-AM pilot provides one concrete existence proof.

References

A. V. Aravindan and M. Kejriwal. Fragile thoughts: How large language models handle chain-of-thought perturbations. arXiv preprint arXiv:2603.03332, 2026.

A. D. Baddeley and G. J. Hitch. Working memory. Psychology of Learning and Motivation, 8:47–89, 1974.

Y. Bai, X. Lv, J. Zhang, H. Lyu, J. Tang, Z. Huang, Z. Du, X. Liu, A. Zeng, L. Hou, et al. LongBench: A bilingual, multitask benchmark for long context understanding. arXiv preprint arXiv:2308.14508, 2024.

S. R. Bowman and G. E. Dahl. What will it take to fix benchmarking in natural language understanding? In Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL), pages 4843–4855, 2021.

D. T. Campbell and D. W. Fiske. Convergent and discriminant validation by the multitrait-multimethod matrix. Psychological Bulletin, 56(2):81–105, 1959.

D. Card, P. Henderson, U. Khandelwal, R. Jia, K. Mahowald, and D. Jurafsky. With little power comes great responsibility. In Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP), pages 9263–9274, 2020.

M. Chen, J. Tworek, H. Jun, Q. Yuan, H. P. de Oliveira Pinto, J. Kaplan, H. Edwards, Y. Burda, N. Joseph, G. Brockman, A. Ray, R. Puri, G. Krueger, M. Petrov, H. Khlaaf, G. Sastry, P. Mishkin, B. Chan, S. Gray, N. Ryder, M. Pavlov, A. Power, L. Kaiser, M. Bavarian, C. Winter, P. Tillet, F. P. Such, D. Cummings, M. Plappert, F. Chantzis, E. Barnes, A. Herbert-Voss, W. H. Guss, A. Nichol, A. Paino, N. Tezak, J. Tang, I. Babuschkin, S. Balaji, S. Jain, W. Saunders, C. Hesse, A. N. Carr, J. Leike, J. Achiam, V. Misra, E. Morikawa, A. Radford, M. Knight, M. Brundage, M. Murati, K. Mayer, P. Welinder, B. McGrew, D. Amodei, S. McCandlish, I. Sutskever, and W. Zaremba. Evaluating large language models trained on code. arXiv preprint arXiv:2107.03374, 2021.

P. Clark, I. Cowhey, O. Etzioni, T. Khot, A. Sabharwal, C. Schoenick, and O. Tafjord. Think you have solved question answering? try ARC, the AI2 reasoning challenge. In arXiv preprint arXiv:1803.05457, 2018.

K. Cobbe, V. Kosaraju, M. Bavarian, M. Chen, H. Jun, L. Kaiser, M. Plappert, J. Tworek, J. Hilton, R. Nakano, C. Hesse, and J. Schulman. Training verifiers to solve math word problems. arXiv preprint arXiv:2110.14168, 2021.

L. J. Cronbach and P. E. Meehl. Construct validity in psychological tests. Psychological Bulletin, 52(4):281–302, 1955.

A. D'Amour, K. Heller, D. Moldovan, B. Adlam, B. Alipanahi, A. Beutel, C. Chen, J. Deaton, J. Eisenstein, M. D. Hoffman, et al. Underspecification presents challenges for credibility in modern machine learning. Journal of Machine Learning Research, 23 (226):1–61, 2022.

K. de Langis, J. I. Park, B. Hu, K. C. Le, A. Schramm, M. C. Mensink, A. Elfenbein, and D. Kang. Strong memory, weak control: An empirical study of executive functioning in llms. arXiv preprint arXiv:2504.02789, 2025.

J. H. Flavell. Metacognition and cognitive monitoring: A new area of cognitive-developmental inquiry. American Psychologist, 34(10):906–911, 1979.

R. Geirhos, J.-H. Jacobsen, C. Michaelis, R. Zemel, W. Brendel, M. Bethge, and F. A.

Wichmann. Shortcut learning in deep neural networks. Nature Machine Intelligence, 2 (11):665–673, 2020.

A. Ghasemabadi and D. Niu. Can llms predict their own failures? self-awareness via internal circuits. arXiv preprint arXiv:2512.20578, 2025.

C. A. E. Goodhart. Problems of monetary management: The U.K. experience. Monetary Theory and Practice, pages 91–121, 1984. Original lecture 1975; source of Goodhart's Law.

C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger. On calibration of modern neural networks. In International Conference on Machine Learning (ICML), 2017.

F. G. Haznitrama, F. R. Ardi, and A. Oh. A neuropsychologically grounded evaluation of llm cognitive abilities, 2026. URL https://arxiv.org/abs/2603.02540.

D. Hendrycks, C. Burns, S. Basart, A. Zou, M. Mazeika, D. Song, and J. Steinhardt. Measuring massive multitask language understanding. In International Conference on Learning Representations (ICLR), 2021a.

D. Hendrycks, C. Burns, S. Kadavath, A. Arora, S. Basart, E. Tang, D. Song, and J. Steinhardt. Measuring mathematical problem solving with the MATH dataset. In Advances in Neural Information Processing Systems (NeurIPS), 2021b.

A. Z. Jacobs and H. Wallach. Measurement and fairness. In Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FAccT), pages 375–385, 2021.

E. Jiang, C. Xu, N. Singh, T. Qiu, and G. Singh. Robust answers, fragile logic: Probing the decoupling hypothesis in llm reasoning, 2026. URL https://arxiv.org/abs/2505.17406.

M. K. Johnson, S. Hashtroudi, and D. S. Lindsay. Source monitoring. Psychological Bulletin, 114(1):3–28, 1993.

N. Kim and S. Schuster. Entity tracking in language models. arXiv preprint arXiv:2305.02363, 2023.

T. Lanham, A. Chen, A. Radhakrishnan, B. Steiner, C. Denison, D. Hernandez, D. Li, E. Durmus, E. Hubinger, J. Kernion, K. Lukosiute, K. Nguyen, N. Cheng, N. Joseph, N. Schiefer, O. Rauber, S. McCandlish, C. Olsson, T. Henighan, T. Hume, Z. Hatfield-Dodds, J. Kaplan, J. Clark, S. R. Bowman, and E. Perez. Measuring faithfulness in chain-of-thought reasoning. arXiv preprint arXiv:2307.13702, 2023.

P. Liang, R. Bommasani, T. Lee, D. Tsipras, D. Soylu, M. Yasunaga, Y. Zhang, D. Narang, et al. Holistic evaluation of language models. Transactions on Machine Learning Research, 2023.

H. Lightman, V. Kosaraju, Y. Burda, H. Edwards, B. Baker, T. Lee, J. Leike, J. Schulman, I. Sutskever, and K. Cobbe. Let's verify step by step. arXiv preprint arXiv:2305.20050, 2023.

X. Liu, H. Yu, H. Zhang, Y. Xu, X. Lei, et al. AgentBench: Evaluating LLMs as agents. In International Conference on Learning Representations, 2023.

D. Manheim and S. Garrabrant. Categorizing variants of Goodhart's law. arXiv preprint arXiv:1803.04585, 2018.

S. Messick. Validity. In R. L. Linn, editor, Educational Measurement, pages 13–103. American Council on Education / Macmillan, 3rd edition, 1989.

G. A. Miller. The magical number seven, plus or minus two: Some limits on our capacity for processing information. Psychological Review, 63(2):81–97, 1956.

A. Miyake, N. P. Friedman, M. J. Emerson, A. H. Witzki, A. Howerter, and T. D. Wager. The unity and diversity of executive functions and their contributions to complex "frontal lobe" tasks. Cognitive Psychology, 41(1):49–100, 2000.

T. O. Nelson and L. Narens. Metamemory: A theoretical framework and new findings. Psychology of Learning and Motivation, 26:125–173, 1990.

M. Nye, A. J. Andreassen, G. Gur-Ari, H. Michalewski, J. Austin, D. Biber, D. Dohan, A. Lewkowycz, M. Bosma, D. Luan, C. Sutton, and A. Odena. Show your work: Scratchpads for intermediate computation with language models. arXiv preprint arXiv:2112.00114, 2021.

H. Putnam. Psychological predicates. In W. H. Capitan and D. D. Merrill, editors, Art,

Mind, and Religion, pages 37–48. University of Pittsburgh Press, 1967.
M. T. Ribeiro, T. Wu, C. Guestrin, and S. Singh. Beyond accuracy: Behavioral testing of NLP models with CheckList. In Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL), pages 4902–4912, 2020.
N. Shinn, F. Cassano, A. Gopinath, K. Narasimhan, and S. Yao. Reflexion: Language agents with verbal reinforcement learning. In Advances in Neural Information Processing Systems, volume 36, 2023.
J. Sweller. Cognitive load during problem solving: Effects on learning. Cognitive Science, 12(2):257–285, 1988.
E. Tulving et al. Episodic and semantic memory. Organization of memory, 1(381-403):1, 1972.
M. Turpin, J. Michael, E. Perez, and S. R. Bowman. Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting. arXiv preprint arXiv:2305.04388, 2023.
J. Uesato, N. Kushman, R. Kumar, F. Song, N. Siegel, L. Wang, A. Creswell, G. Irving, and I. Higgins. Solving math word problems with process- and outcome-based feedback. arXiv preprint arXiv:2211.14275, 2022.
C. Wang and J. V. Sun. Unable to forget: Proactive interference reveals working memory limits in llms beyond context length, 2025. URL https://arxiv.org/abs/2506.08184.
W. Wang, J. Min, and W. Zou. Intelligence degradation in long-context llms: Critical threshold determination via natural length distribution analysis. arXiv preprint arXiv:2601.15300, 2026.
S. Xie, Z. Shi, H. Shen, G. Huang, Y. Ma, and X. Jing. M3-bench: Process-aware evaluation of llm agents social behaviors in mixed-motive games. arXiv preprint arXiv:2601.08462, 2026.
Z. Yang, P. Qi, S. Zhang, Y. Bengio, W. W. Cohen, R. Salakhutdinov, and C. D. Manning. HotpotQA: A dataset for diverse, explainable multi-hop question answering. In Proceedings of EMNLP 2018, 2018.
S. Yao, H. Chen, J. Yang, and K. Narasimhan. WebShop: Towards scalable real-world web interaction with grounded language agents. In Advances in Neural Information Processing Systems (NeurIPS), 2022.
R. Zellers, A. Holtzman, Y. Bisk, A. Farhadi, and Y. Choi. HellaSwag: Can a machine really finish your sentence? In Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL), 2019.

A Full Related Work Table

Table 7. Representative benchmarks and evaluation work, grouped by measurement type and CEF gap addressed.

Work	Measures	CEF gap addressed
MMLU [Hendrycks et al., 2021a], ARC [Clark et al., 2018], HellaSwag [Zellers et al., 2019]	Knowledge recall	No memory structure, no process
HumanEval [Chen et al., 2021], MATH [Hendrycks et al., 2021b], GSM8K [Cobbe et al., 2021]	Output correctness	No reasoning fidelity
HotpotQA [Yang et al., 2018], WebShop [Yao et al., 2022]	Task completion	No metacognition, no memory typ
AgentBench [Liu et al., 2023]	Sequential task success	No within-task process
Reflexion [Shinn et al., 2023]	Post-reflection improvement	Monitoring $\neq$ control (conflated)
PAE [Xie et al., 2026]	Corruption rate	Motivates CEF; no per-dimension measurement
NeuroCognition [Haznitrama et al., 2026]	Cognitive ability (IQ)	Ability $\neq$ process quality; no load variation
Strong Mem. & Weak Ctrl. [de Langis et al., 2025]	Executive function	Isolated tasks; no multi-turn load
Entity Tracking [Kim and Schuster, 2023]	Entity-state	Prior art for WMF-AM
Intell. Degradation [Wang et al., 2026]	Load degradation	Prior art for CLA
HELM [Liang et al., 2023]	Behavioral outcomes	No cognitive processes
Long-context [Bai et al., 2024]	Retrieval from static window	WMF-AM targets active transformation
ECE calibration [Guo et al., 2017]	Confidence scalar	MCC decomposes into monitoring/control

Table 8. Development status of all 11 CEF sub-dimensions. Status key—A: robust differentiation + multi-seed reliability; B: differentiation shown, limited seeds/models; C: probe redesigned, initial results; D: ceiling/floor effect, needs redesign; E: protocol specified, no pilot data.

Sub-dim	Status	Evidence	Limitation
WMF-AM	A^+	Δ=0.933, 4-seed all 20, 3-template (τ=0.524), yoked control, predictive τ=0.612	All open-weight
WMF-IM	B	spread 0.480, N=5	Single paradigm (hard)
WMF-IR	D	Ceiling (= 1.0 all models)	Needs redesign
MCC-MA	B	4/7 models differentiated	Floor for 3 models
MCC-CE	C	v2: error 23–40%, near-zero flagging (N=8)	Binary protocol
EMC-lite	B	Source-conflict spread 0.400	Single seed
CLA-DC	$B^\dagger$	Correlates τ=0.905 with WMF-AM	†Provisionally collapsed onto WMF
CLA-CR	C	v3 spread 1.171, N=7	v1–v2 failed; 2 seeds
CLA-RA	E	—	Awaits full study
EMC-TO	E	—	Awaits full study
EMC-EI	E	—	Awaits full study

Table 9. CEF at a glance. Each sub-dimension's task type, scoring rule, and pilot status. Tier 1 = novel construct; Tier 2 = novel reframing. Pilot: ✓ = pilot-supported, ∼ = preliminary, — = protocol only.

Sub-dim	Task (one-line)	Scoring	Tier	Pilot
WMF-AM	Track K sequential state updates, report final state	Exact match	1	✓
WMF-IM	Recall N facts amid M distractors	Recall accuracy	2	∼
WMF-IR	Recall List-A after List-B interference	Proactive interference δ	2	—
MCC-MA	Predict which of 20 answers are wrong	Pearson r (predicted vs. actual)	1	∼
MCC-CE	Flag uncertain answers, revise them	Correction efficacy	1	∼
EMC-TO	Reorder non-chronological events	Kendall τ vs. ground truth	1	—
EMC-EI	Discriminate similar episodes	Accuracy under similarity	1	—
EMC-SA	Attribute facts to source + episode	Joint attribution accuracy	2	—
CLA-DC	Performance across 5 load levels	Smoothness (SD of Δs)	1	—
CLA-RA	Response elaboration vs. difficulty	Pearson r	1	—
CLA-CR	Easy-task recovery after high-load block	Recovery rate	1	∼

B Development Status and Summary Tables

C Glossary

D On Cognitive-Science Labels

We use labels like "working memory" and "metacognition" as operational shorthand for specific behavioral probes, not as claims about LLM internal mechanisms. "WMF-AM" means: "accuracy on sequential state-tracking tasks degrades with depth K"; it does not mean LLMs have a working memory buffer. "MCC-MA" means: "accuracy of self-error prediction on a 20-item probe"; it does not mean LLMs have phenomenal metacognition. The labels are justified because the probes are adapted from validated human measurement paradigms, and they organize what would otherwise be an unstructured collection of tests.

Two caveats are essential. First, we do not require LLMs to have phenomenal cognition or internal states analogous to human working memory. CEF dimensions are process-sensitive behavioral probes that borrow measurement logic from cognitive science. Second, CEF measurements are behavioral (prompt/response), designed to be differentially sensitive to process quality: high MCC-MA requires accurate self-error prediction; graceful CLA-DC requires proportional rather than catastrophic degradation.

E Full Sub-dimension Definitions

WMF-AM (Active Manipulation, Tier 1). Uses a complex span task: initial state, K sequential modifications, report final state. Depths $K \in \{2, 4, 6, 8, 12\}$, three surface forms. Prior entity tracking [Kim and Schuster, 2023] tests passive retrieval; WMF-AM tests sequential transformation.

WMF-IM (Information Maintenance, Tier 2). Embeds N facts in a passage with M distractors; measures recall accuracy. Extends Wang and Sun [Wang and Sun, 2025] to natural-language multi-turn contexts.

Table 10. Glossary of key terms and acronyms.

Term	Definition
CEF	Cognitive Evaluation Framework (this paper)
WMF-AM/IM/IR	Working Memory Fidelity: Active Manipulation / Information Maintenance / Interference Resistance
MCC-MA/CE	Metacognitive Calibration: Monitoring Accuracy / Control Efficacy
EMC-TO/EI/SA	Episodic Memory Coherence: Temporal Ordering / Episodic Interference / Source Attribution
CLA-DC/RA/CR	Cognitive Load Adaptation: Degradation Curve / Resource Allocation / Capacity Recovery
AGENT-PQ/TC	Agent Process Quality / Task Completion (external evaluation)
CONV-RFPOC	Convergent validity probe: Reasoning Fidelity–Product of Cognition

WMF-IR (Interference Resistance, Tier 2). Adapts the List-A/List-B proactive interference paradigm. Currently shows ceiling (= 1.0 all models); needs redesign.

MCC-MA (Monitoring Accuracy, Tier 1). Model predicts which of 20 answers are incorrect; Pearson r measures monitoring accuracy. Unlike Gnosis [Ghasemabadi and Niu, 2025] (requires internals), MCC-MA uses prompting under multi-turn load.

MCC-CE (Control Efficacy, Tier 1). Model flags uncertain answers and revises them. Decomposes into flagging rate, correction efficacy, and false alarm rate.

EMC-TO/EI/SA. Temporal ordering (Kendall τ), cross-episode discrimination under similarity, and joint source-episode attribution. All proposed, no pilot data.

CLA-DC (Degradation Curve, Tier 1). Smoothness score (SD of successive Δs). Provisionally collapsed onto WMF-AM ($\tau = 0.905$).

CLA-RA/CR. Resource allocation (difficulty–elaboration correlation) and capacity recovery (post-load accuracy recovery). CLA-RA has no pilot data; CLA-CR-v3 spread = 1.171.

Composites (all exploratory). WMF = $0.50\times$AM$+0.30\times$IM$+0.20\times$IR; MCC = $0.55\times$MA$+0.45\times$CE; EMC = $0.40\times$EI$+0.35\times$TO$+0.25\times$SA; CLA = $0.50\times$DC$+0.30\times$RA$+0.20\times$CR. We do not propose an overall CEF composite; per-dimension profiles are the primary unit of analysis.

F Construct-to-Probe Mapping and Falsifiability

Table 11. Construct map for pilot-tested probes: manipulated factor, observable signature, rival explanations, and falsification criteria.

Probe	Manipulated factor	Observable signature	Plausible rivals	Expected pattern	Falsified if
WMF-AM	Depth K (number of sequential state updates)	Accuracy × depth interaction	Instruction following; tokenizer parsing; template familiarity; long-context retrieval	Low τ with completion; monotonic depth decline; $\tau_{\text{bare,chat}} > 0.4$; yoked control $\gg$ WMF-AM	High τ with completion ($\geq$ 0.55); no depth effect; template-specific rankings; control degrades at same rate
MCC-MA	Error-laden vs correct items	Pearson r (predicted vs actual errors)	Calibration shortcuts; always-predict-correct bias	Low τ with WMF-AM; model-specific prediction	Uniform predictions; $\tau \approx 1$ with WMF-AM
EMC-lite	Conflict vs control episodes	Source attribution accuracy under interference	Recency bias; keyword matching	Low τ with WMF-AM; fault > control difficulty	No fault–control difference; $\tau \approx 1$ with WMF-AM

G Rival Explanations

For each probe, plausible alternative explanations remain open. WMF-AM: (a) Instruction following—models may fail at depth because they lose track of instructions. The depth gradient partly addresses this. (b) Long-context retrieval—the control addresses this for 5/7 models. (c) Template artifacts—cross-template stability (τ = 0.685–0.781) mitigates but does not eliminate this. MCC-MA/CE: (a) Calibration shortcuts. (b) Prompt sensitivity. EMC-lite: Recency bias; keyword matching. These are the primary targets for the $N \geq 15$ study.

H Study Matrix

Table 12. Study matrix: all experiments, models, and sample sizes. †llama3.1:70b added for MCC-CE-v2 only.

Experiment	Metric	N_{models}	Seeds	Trials/model
Outcome Correctness (v2)	Accuracy	20	1	100
WMF-AM (dissociation)	Exact match	20	4	180
Agent Validation	Process quality	20	1	10
Template Harmonization	Rank stability	15	1	135
Yoked Cancellation Control	Exact match	15	1	100
MCC-MA	Pearson r	15	1	20
MCC-CE-v2	Flagging rate	8†	1	50
EMC-lite	Composite	7	1	18
CLA-CR-v3	Recovery rate	7	2	7 blocks
AGENT-PQ/TC	Rubric / binary	7	1	10
CONV-RFPOC	Fidelity	7	1	15

I Yoked Cancellation Control Details

The yoked cancellation control uses operation pairs that cancel: "Alice gains 3 points" followed immediately by "Alice loses 3 points." The prompt format, entity names, and depth structure match WMF-AM exactly, so the arithmetic parsing demand is identical. The correct answer equals the initial state, but unlike the original inert control, the model must process each operation to verify cancellation. This eliminates the trivial shortcut of ignoring operations.

J Measurement Robustness

Multi-seed reliability. All 15 models evaluated with 4 independent random seeds × 15 probes. Original-model cross-seed $\tau = 0.798$. Expansion-model SDs: phi3:14b = 0.094, gemma2:9b = 0.144, qwen2.5:3b = 0.122, llama3.2:3b = 0.122, deepseek-r1:7b = 0.114, mixtral:8x7b = 0.116, command-r:35b = 0.100, yi:34b = 0.084 (mean SD = 0.108).

Cross-template stability. Phase 5 administered WMF-AM under three prompt wrappers (bare, chat, chain-of-thought) across all 15 models. Bare ↔ chat: $\tau = 0.524$ ($p = 0.006$)—significant rank preservation. Bare ↔ CoT: $\tau = 0.105$ ($p = 0.627$)—CoT pushes all models to ceiling (≥ 0.93), eliminating variance. Chat ↔ CoT: $\tau = 0.314$ ($p = 0.114$). The CoT ceiling effect is interpretively meaningful: externalized reasoning via scratchpad bypasses the working-memory bottleneck that WMF-AM is designed to measure.

Non-ceiling completion proxies. τ(WMF-AM, MMLU) = 0.356 ($p = 0.317$; range 0.25–0.72); τ(WMF-AM, GSM8K) = 0.150 ($p = 0.645$; range 0.05–0.83). Scores from published model cards, not re-evaluated under pilot setup.

Depth profile.

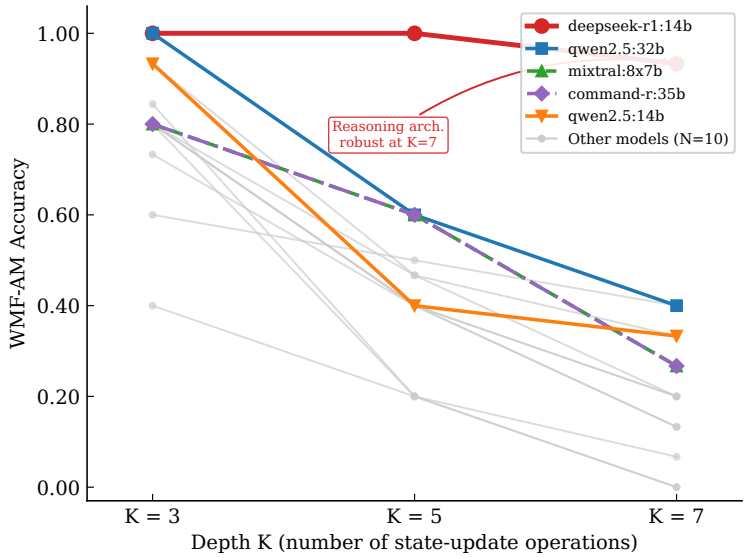

Figure 4. WMF-AM accuracy by depth K (multi-seed means, N=7). deepseek-r1:14b maintains near-perfect accuracy; others degrade at K=7.

Model	F-Det	F-Src	F-Rec	F-Comp	C-Rec	ΔRec
qwen2.5:14b	0.917	0.750	0.917	0.917	0.972	−0.056
qwen2.5:32b	0.917	0.750	0.917	0.917	1.000	−0.083
deepseek-r1:14b	0.889	0.917	0.917	0.912	1.000	−0.083
llama3.1:8b	0.778	0.583	0.750	0.771	0.889	−0.139
qwen2.5:7b	0.861	0.583	0.639	0.736	0.722	−0.083
mistral:7b	0.750	0.417	0.639	0.685	0.722	−0.083
gemma2:27b	0.861	0.583	0.528	0.667	0.611	−0.083
Δ	0.167	0.500	0.389	0.250	0.389	—

Table 13. EMC-lite (N=7, hardened protocol). Source attribution shows strongest differentiation (spread = 0.500).

K EMC-lite Results

EMC-lite provides preliminary support for the source-conflict slice. It does not establish EMC as a whole; single-seed evidence only. τ(WMF-AM, EMC-lite) = 0.250: substantial independence.

L MCC-CE-v2 Results

The v1 floor was an artifact of easy questions; v2 models make many errors but show near-zero self-flagging under the binary protocol. Alternative designs (continuous confidence, forced ranking) might reveal partial capacity.

M Negative OOD Result

Two OOD stress-test designs (v1: 5-step NL tasks; v2: K=8–9 interference tasks) produced anti-correlated CEF–OOD relationships across model pairs. We make no predictive claim. Full assessment requires $N \geq 11$ with process-sensitive downstream tasks.

N Full Validity Analyses

The validity summary table appears in the main text (Table 6). Below we provide additional details.

Table 14. MCC-CE-v2 (N=8). Error rates 23–40% (floor eliminated); flagging near zero for all models.

Model	Balanced (n=30) Error%	Balanced (n=30) Flag%	Trap-only (n=20) Error%	Trap-only (n=20) Flag%
qwen2.5:7b	26.7	0.0	40.0	0.0
qwen2.5:14b	36.7	0.0	40.0	0.0
qwen2.5:32b	30.0	0.0	25.0	0.0
llama3.1:8b	26.7	0.0	40.0	0.0
llama3.1:70b	—	—	25.0	0.0
gemma2:27b	30.0	11.1	30.0	0.0
deepseek-r1:14b	23.3	0.0	35.0	0.0
mistral:7b	30.0	0.0	40.0	0.0

Convergent validity details. AGENT-PQ: 10 multi-step scenarios scored by LLM judge (GPT-4o) on 4 rubric dimensions. CONV-RFPOC: 15 reasoning problems with counterfactual chain perturbation. RFPOC $= 1 - P(\text{same answer} \mid \text{perturbed chain})$. Full protocol in Appendix O. Limitation: LLM judge may introduce shared bias.

Construct separation (WMF-AM $\leftrightarrow$ CLA-DC). $\tau = 0.905$: no evidence of separability at N=7. Both involve state tracking under load. Planned discriminant test at $N \geq 15$. CLA-CR-v3 provides a structurally independent CLA measure. Even if CLA-DC merges with WMF-AM, WMF-AM, MCC-MA, and EMC-lite remain separable ($\tau \leq 0.250$).

Sensitivity. LOMO for CONV-RFPOC vs AGENT-PQ: τ ranges 0.600–1.000 (mean 0.833). Leave-one-family-out (Qwen, N=4): $\tau = 0.667$. After Holm correction, RFPOC–AGENT-PQ remains significant ($p_{\text{adj}} = 0.021$).

Multiple testing. Three primary hypotheses; all others exploratory.

Benchmark governance. All headline results use frozen protocol versions. Probe revisions were motivated by floor/ceiling effects, not model-specific tuning.

O AGENT-PQ Evaluation Protocol

AGENT-PQ uses 10 multi-step scenarios where each model operates within a tool-use scaffold. A separate LLM judge (GPT-4o, temperature = 0) scores each trajectory on four rubric dimensions: plan coherence, state tracking, error recovery, goal adherence. AGENT-PQ = mean of 4 rubric scores/4. AGENT-TC is the binary success rate on the same scenarios.

CONV-RFPOC administers 15 reasoning problems with counterfactual chain perturbation. RFPOC $= 1 - P(\text{same answer} \mid \text{perturbed chain})$.

P Complete Experimental Protocols

Full prompt templates, parameter specifications, scoring procedures, and reference implementations for all 11 sub-dimensions are provided in the supplementary materials.

Q Weight Sensitivity Analysis

Three weighting schemes (proposed, equal, PCA-derived) are evaluated. Per-dimension profiles are primary; composites are exploratory only.

R Dissociation Case Studies

Matched dissociation pairs: deepseek-r1:14b and qwen2.5:14b share similar EMC-lite (0.912 vs. 0.917) but differ on WMF-AM (0.983 vs. 0.467). mistral:7b and llama3.1:8b have similar WMF-AM but differ on EMC-lite. These cross-dimension pairs anchor the multiple realizability argument.

S Worked Example: WMF-AM at Depth K=3

Setup. "A warehouse has 3 items: Widget=50, Gadget=30, Sprocket=20." Three updates: "Ship 10 Widgets," "Receive 25 Gadgets," "Transfer 5 Sprockets to overflow."
Correct. Widget=40, Gadget=55, Sprocket=15.
Scoring. Exact match per item. deepseek-r1:14b scores 1.000 at K=7; llama3.1:8b scores 0.050. Both score 0.90–0.95 on outcome correctness. This is the Completion Fallacy in action.

T Novelty Classification

Table 15. Novelty tier for all 11 sub-dimensions. Tier 1 = novel; Tier 2 = novel reframing.

Sub-dim	Tier	Prior art / differentiation
WMF-AM	1	Entity Tracking [Kim and Schuster, 2023]: passive retrieval; WMF-AM = sequential transformation
WMF-IM	2	Wang & Sun [Wang and Sun, 2025]: key-value pairs; WMF-IM = NL multi-turn
WMF-IR	2	Extended to domain-varied materials
MCC-MA	1	Gnosis [Ghasemabadi and Niu, 2025] requires internals; MCC-MA = prompting only
MCC-CE	1	Reflexion [Shinn et al., 2023]: aggregate; MCC-CE decomposes detection/correction
EMC-TO	1	No prior LLM temporal ordering measurement
EMC-EI	1	No prior cross-episode interference diagnostic
EMC-SA	2	Interference-conditioned attribution
CLA-DC	1	Intell. Degradation [Wang et al., 2026]: magnitude; CLA-DC = shape
CLA-RA	1	Not previously measured
CLA-CR	1	Not previously measured